

# The System (RDBMS)

## 1. Appliance

*The Database System as an Appliance (a review)*

## 2. Transaction Theory

## 3. Transaction Management

## 4. Other RDBMS Functions

## 5. Application Programming



# Version

<b>initial</b>	<i>2020-04-07</i>	Parke Godfrey
<b>adapted</b>	<i>2020-11-16</i>	Wenxiao Fu
<b>last revised</b>	<i>2020-11-17</i>	Parke Gdofrey



# Acknowledgments

## Thanks to

- *Jeffrey D. Ullman*  
for slides related to the textbook (*Stanford slides*)
- *Jarek Szlichta*  
for significant refinements to the *Stanford slides*
- *Wenxiao Fu*  
for significant refinements



## Printable version of talk

- Follow this link:  
**The System [to pdf]**.
- Then *print* to get a **PDF**.

*(This only works correctly in **Chrome** or **Chromium**.)*



# 1. Appliance

## *The Database System as an Appliance (a review)*



# The RDBMS as an Appliance

This course has been all about *how to use* an **RDBMS** — a *relational database management system* — as an *appliance* to help

- *manage* and
- *access*

our data.



## Design (& Schema)

The *first* main section of the course was on *design*: how to map out a logical organization for the data we want to keep, called a *schema*.

- The **RDBMS** does not help us to do this design,
- but it can manage the *schema* and the data put into it for us.

The *relational model* does provide us a schema “language” that dictates *how* we can organize our data.



## Queries (in SQL)

The *second* main section of the course was on *how* to access, and manipulate, the data in our database.

- The paradigm for this is to *query* over the database.
  - The nearly universal query language for this — for relational database systems — is SQL.
- 

What does the **RDBMS** provide?

- It does not help us to *compose* our SQL queries; only we know what we want.
- But it does evaluate our queries against the database — the schema and the data it contains — efficiently.





# The RDBMS as a Blackbox

SQL, our way of interacting with our data, is *declarative*.

- We compose a query to state logically *what* we want, but we do not spell out a procedure for *how* this is to be accomplished.

The **RDBMS** has to do the *how*.

- This is a *hallmark* of **RDBMS**'s, and what has made them so widely useful and used.

As such, in *this* course, we are *clients* of the **RDBMS**. We have to know how to work with it, but we do not have to know *how* it works.

We can treat the **RDBMS** as a *blackbox*.



## ***What else should RDBMS clients (users) know?***

What other functions of the **RDBMS** should users (**us**) master to use it effectively?

And that we should cover in this course?

*Many!* Too many to cover in a course of 12 to 15 weeks.

...but there are some important topics still to cover.

And, to be *expert* **RDBMS** users, do we need to look “under the hood” of the system at all?

Doesn't relational's *declarative* ideal mean we do not need to?!



## Changing Data (*data in motion*)

We have mostly looked at the database as *static*, so far.

- We made the schema, then filled it with some data.
- We then obsessed on how we query over the data to retrieve the information we want.
- But data will be added and updated all the time!

So we also need a way to *manage* these *transactions*.

- *Integrity*: No transaction should be able to violate the schema's rules!
- *Concurrency*: Many transactions can be occurring “at the same time”.



# Concurrency

An **RDBMS** lets a database be shared by many, many clients at the same time.

- The *data* in the database is usually being added to, and modified, quite rapidly.
- And this is often being done *at the same time* by many clients.



# Application Programming

## *and learning to live with others (concurrently)*

As database users, we need to understand the consequences of “competing” with others for use of the database to be able to use the **RDBMS** effectively.

---

We need to know the *logic* of *how* the **RDBMS** deals with our “transactions”.

Oh, and to write sophisticated applications that use the **RDBMS** to manage the data, we cannot live with *just* SQL.

How we can couple SQL with *programming languages* (e.g., **Java**) is needed.

This is called *application programming*.



# Application Programming Interface(s)

Before we can dive into the “API” support **RDBMS**'s provide for application programming, we need to understand *how* the database system handles our transactions against the database.

---

This motivates the two key topics that we, as expert database users, need to master under this *third* section of the course:

1. transactions, *theory & management*
2. the application-programming interfaces



## 2. Transaction Theory



# Important System Aspects

- For permanence, data is on disk.
- To work on data, it must be in main memory. (But main memory is volatile!)
- Main memory is orders of magnitude faster than secondary (“disk”) memory.





# Primitive Operations

- *Read* a piece of data.
- *Write* a piece of data.

System-wise, a *transaction* is just a sequence of reads and writes.



# Transactions

A *transaction* is a sequence of connected actions — *querying* and *updating* data — over the database.

---

The **RDBMS** ensures that our schema's rules are never violated.

It still needs to do this when there are transactions being processed at the same time (*concurrently*) and the transactions might logically conflict.

*How* transactions can logically conflict and *what* guarantees the **RDBMS** can provide to prohibit possible logical conflicts is called *transaction theory*.

And *how* the **RDBMS** performs this is called *transaction management*.



# Transaction Theory

What can go “logically” wrong when two or more transactions “conflict” is *not* at all obvious.

---

What *can* go wrong was enumerated by **Jim Gray** (**@ Wikipedia**), along with *what* transactional *properties* would prevent these issues and *how* to guarantee the properties (*transaction management*).

Dr. Gray won the *Turing Award* in 1998 for this work, “for seminal contributions to database and transaction processing research and technical leadership in system implementation.” (**Jim Gray** (**@ Wikipedia**))

Tragically, Dr. Gray vanished at sea piloting his 40-foot sailboat near San Francisco in *January 2007*.



# Anomalies

## vs the Transactional Properties

There are a number of logical *anomalies* that can arise when transactions “conflict”.

This is akin to our study of *data anomalies* that could arise when a schema is not *normalized*.

---

A set of guiding principles — *transactional properties* — if adhered to will guarantee those anomalies cannot happen.

Again, analogous to how refining a schema to be in BCNF guarantees many types of *data anomalies* cannot occur within the schema.

Which to study first? In this case, the *properties* give us good insight. So we will start with those. (Of course, Dr. Gray had to arrive at the properties from understanding the anomalies!)



# The ACID Properties

- **A:** *atomicity*
  - **C:** *consistency*
  - **I:** *isolation*
  - **D:** *durability*
- 

These are the *four* properties we want the **RDBMS** to guarantee for us via its *transaction management*.

Each property says something about how any given transaction “sees” the database.

Recall, the data in the database is changing over time.



# Atomicity

A transaction is a sequence of *actions* against the database (perhaps coming from an application program or a user with a “PostgreSQL” shell open).

---

Atomicity is also referred to as the *all-or-nothing* principle.

A number of the transaction's actions might make changes to the database. When the transaction is *done*, either *all* of the transaction's changes should now be part of the database (for any future transactions to see), or *none* should be (as if the transaction never existed)!

This is called *atomic* because it looks like the *transaction* was one large “action” that either happens or doesn't.



## Atomicity: Commit or Rollback

*Atomicity* means that a transaction has one of two fates: *all* of its changes were *committed* to the database so now others can see the changes; or none were, as if the transaction *never* existed.

---

In the first case, we say the transaction *committed*. In the second case, we say the transaction was *rolled back*.

In fact, there are two *special* actions a transaction can make: **commit** and **rollback**.

The database system *tries* to honour a **commit**, but cannot always. If it cannot, the transaction is *rolled back*.

And a transaction can *vanish* itself by calling **rollback**!

(The action **rollback** has the synonym **abort**.)



# The Idea of Rollback

This notion that a transaction can be “undone” is very powerful!

We shall see it is a necessary consequence of **ACID**.

- We have to accept that our transaction might be *rolled back* by the database system at times, no fault of ours.
- But we can also “abort the mission” our transaction if we need to by *rolling back* ourselves.





# Triggers & Rollback

Can a *trigger* be used as an *integrity constraint* to prohibit certain types of changes to be made?

*Well, yes!*

---

For example, let us say we want to prohibit that any employee's salary could be updated to a lower value.

- An *update* trigger can monitor for updates to *salary* in **Employee**. (the *event*)
- It checks whether the old salary is more than the new salary. (the *condition*)
- And what is its *action*? **rollback!**

The trigger's actions are within the scope of the transaction that *triggered* it.



# Consistency

The **RDBMS** does not allow any change to the data that would violate the schema.

We call this *consistency*.

Therefore, no transaction can ever *commit* that would leave the data of the database in a state that violates the schema's constraints.

This is something we are already quite used to.



# Isolation

Transactions could not conflict if they were completed one at a time: all my transaction's changes are made and committed, then the next person's, etc.

---

Why not limit the **RDBMS** to handle *one* transaction at a time?

- It would be highly inefficient! A database is shared.
- You might have to wait a long, long time connecting to the database to do something.



## Isolation: the illusion

So we need to let transactions run concurrently, for efficiency's sake.

---

But changes to the database should be *logically* equivalent to *some* order of the transactions happening as if one at a time, sequentially.

We want to provide the *illusion* that the transactions are *isolated*, happening one at a time.



# Transaction Schedules

A *schedule* is the *sequence* of actions executed by the database system.

- The server machine can only do one action at a time. (So no true *parallelism* here.)
- The actions from different transactions may be *interleaved* in the schedule (giving us *concurrency*).

If transactions's actions are not interleaved at all in the schedule, we call the schedule *serial*. This is *true* isolation.

When a schedule's results on the database are *equivalent* to that of some *serial* schedule, we call the schedule *serializable*.

Serializable schedules guarantee the property of *isolation*, the illusion of *true* isolation.



# The Snapshot Principle

Because a transaction (X-act) sees the database as it looks as *after* the X-acts that *logically* came *before* it and sees *none* of the changes by X-acts that logically come *after* it (with respect to the equivalent serial schedule), it is as if the X-act sees the database as it is at *a single point in time* (the *snapshot*).

*Isolation* is synonymous with the *snapshot* principle in the same way that *atomicity* is synonymous with the *all-or-nothing* principle.



# Durability

*Durability* simply states that once a X-act has been *committed* by the database system, the changes made by the X-act are now visible by any X-act that comes after.

**Commit** means the changes are now permanently part of the data. (Well, at least until someone else comes along and changes the same data again!)



# Serializable

We use the term *serializable* for the highest setting of *isolation level* for transactions. *Serializable* guarantees the **ACID** properties are never violated.

---

**ACID** — well, *atomicity* and *isolation* — guarantees that logical **anomalies** cannot occur, where we see “impossibilities” or *phantom* data that is not really there.

There are steps down in *isolation level* that trade off allowing the possibility of some types of anomalies for increased concurrency. (Like dropping down in normal forms.)





## Example: Interacting Processes

bar	beer	price
Joe's Bar	Bud	2.50
Joe's Bar	Miller	3.00

- Sally is querying **Sells** (above) for
  - Q1. the highest price, and
  - Q2. the lowest price Joe's Bar charges.
- Joe decides to
  - Q3. stop selling Bud and Miller (deletion), but
  - Q4. to sell only Heineken at \$3.50 (insertion).



# Sally's Program

Sally executes the following two SQL statements called (max) and (min) to help us remember what they do.

Q1. (*max*)

```
SELECT MAX(price) FROM Sells  
WHERE bar = 'Joe''s Bar';
```

Q2. (*min*)

```
SELECT MIN(price) FROM Sells  
WHERE bar = 'Joe''s Bar';
```



## Joe's Program

At about the same time, Joe executes the following steps: *(del)* and *(ins)*.

Q3. *(del)*

```
DELETE FROM Sells  
WHERE bar = 'Joe''s Bar';
```

Q4. *(ins)*

```
INSERT INTO Sells VALUES  
( 'Joe''s Bar', 'Heineken', 3.50 );
```



# Interleaving of Statements

Except

- *(max)* must come before *(min)*, and
- *(del)* must come before *(ins)*,

there are no other constraints on the order of these statements, unless we group Sally's and/or Joe's statements into transactions!



## Example: Strange Interleaving

Suppose the steps execute in the order  $(max)(del)(ins)(min)$ .

What the highest price and lowest price does Sally get?

Joe's Prices:	{2.50,3.00}	{2.50,3.00}	{3.50}
Statement:	(max)	(del)	(ins) (min)
Result:	3.00		3.50

Sally sees  $MAX < MIN$ !



## Fixing the Problem by Using Transactions

If we group Sally's statements (max)(min) into one transaction, then she cannot see this inconsistency.

She sees Joe's prices at some fixed time. Either

- before (*del*),
- after (*del*) but before (*ins*), or
- after (*ins*),

but the MAX and MIN are computed from the same prices.



## Another Problem: Rollback

Sally	Joe
	<i>(del)</i>
	<i>(ins)</i>
<i>(max)</i>	
<i>(min)</i>	
commit	
	rollback

- If Sally executes her statements after *(ins)* but before the rollback,
- she sees a value, 3.50, that never existed in the database.



# Solution

If Joe executes (del)(ins) as a transaction, its effect cannot be seen by others until the transaction executes COMMIT.

- If the transaction executes ROLLBACK instead, then its effects can never be seen.





## Anomalies: Unrepeatable Reads / RW conflicts

Sally	Joe
<i>(max)</i>	
	<i>(del)</i>
	<i>(ins)</i>
	commit
<i>(max)</i>	
commit	

- There are two different highest prices with the same beer!



# SLIDEDECK in progress...



# Application Programming

Any time a *program* connects to an **RDBMS** to issue SQL statements against a database, that is a *transaction*.

But how do we write programs that can “talk” to a database, “transact” with it?

Most interaction with database systems is via programs, and not through direct SQL from a user.

Fortuantely, how to *couple* programs to interact with a database is very standardized, just like SQL.



## CLI (e.g., JDBC)

We need a *library* for our programming language that provided an API of calls so we can set up interaction with a database system

- This is generically called the *call library interface*)
- For Java, the “CLI” is called JDBC, for *Java-Database Connectivity*.

(I will focus on JDBC here. But CLI's for other languages are fundamentally the same.)

The *only* way to interact with a database under an **RDBMS** is through SQL.

Therefore, JDBC must provide a way for a program to *send* SQL to the database server and to get the results back.



## The JDBC Exchange

1. *sql-str*: In the **Java** program, we can store or build an SQL statement in a string.
2. *conn*: A call establishes a connection to a database server for a given database.
3. *statement*: Using the *conn* object, the *sql-str* is *prepared*. This sends the SQL to the database server, which parses it and “prepares” to execute the SQL commands when called.
4. *result-set*: calling *execute* on *statement* invokes the database server to run the SQL statement and return the results.



## The Result Set

Let us consider an SQL query that was *prepared* and *executed*. How does the **Java** program process the results? A query returns a *set* — okay, a *multi-set* — of tuples.

It is done via a specific type of *iterator*, a *cursor* (class `ResultSet` in JDBC). With the iterator, the program can iterate through the tuples of the answer set.



# The Cursor

On the database-side of the fence, a *cursor* is an open *transaction*. And the snapshot semantics is ensured.

---

See **JDBC: A small tutorial** [**pdf**] to see the concrete steps.

And see **JDBC: Example** for an example.



# Conclusion

And that concludes this review of *The System* for *transactions* and *applications*.

---

Study well!

Keep well!

The System | EECS-3421A: *Introduction to Database Systems* | Fall 2020

